

APPENDIX: PROOF OF CORRECTNESS

PROPOSITION 1. (*Uniform integrity*) *For any message m , every process p delivers m at most once, and only if p is a destination of m and m was previously multicast.*

PROOF: Process p delivers m at Task 6 if m 's state is ORDERED. After delivering m , p sets m 's state to DONE, and thus m cannot be delivered more than once.

Let c be the client that multicasts m to groups in dst , and let p be in group g . From Task 6, p only delivers m if it is in p 's M buffer and m 's state is ORDERED. Message m 's state is set to ORDERED in Task 5 if its current state is MCAST. A message's state is set to MCAST in procedure *Relay*, which is invoked in two cases: (a) by client c upon multicasting m (Task 1) to groups in dst , in which case $g \in dst$; or (b) by some process q that suspects c (Task 10), has m in its buffer in state MCAST, and g is a destination of m . In case (b), m was written in q 's buffer either (b.1) directly by c or (b.2) indirectly by some other process. In any case, there is some process r such that m is included in r 's buffer by c . It follows from Task 1 that p is a destination of m and m was multicast by client c . \square

LEMMA 1. *If all correct processes in the destination of an atomically multicast message m have m in their M buffer in the MCAST state, then they eventually set m to the ORDERED state.*

PROOF: Let m be addressed to groups in dst and q be a correct process addressed by m . We claim that for each $h \in dst$, q will have a timestamp for h that is acknowledged by a quorum of processes in h . By the leader election oracle and the fact that each group has a majority of correct processes, group h eventually has a stable correct leader l . Either (a) l executes Task 2 and proposes its clock value as h 's timestamp or (b) l executes Task 7 to replace a suspected leader. In (b), l sends a CATCH_UP message to all processes and will receive for each group $g \in dst$ the timestamp proposed in g , if any, and the corresponding acknowledgements from processes in g (Task 8). For the case where $h = g$, l will pick the timestamp decided by a previous leader or choose one if no timestamp has been decided (Task 9). Thus, in both cases (a) and (b), the leader writes the chosen timestamp in the M buffer of each process in h and in the leaders of other groups in dst . From Task 3, every follower in h will acknowledge this timestamp in the buffer of each process in the destination of m . From Task 4, when l has a timestamp from $g \neq h$, l writes the timestamp in the buffer of its followers, which concludes the claim. Therefore, eventually q has a timestamp for every group in dst , can compute m 's final timestamp, and set m 's state as ORDERED.

LEMMA 2. *If a correct process p has an atomically multicast message m in its M buffer in the ORDERED state, then p eventually delivers m .*

PROOF: Assume for a contradiction that q does not deliver m . Thus, there is some message m' in the buffer such that $m \neq m'$, m' 's timestamp is smaller than m 's timestamp, and m' 's state is not DONE.

We first show that any message added in the buffer after m becomes ORDERED has a timestamp bigger than m 's timestamp. Message m only becomes ordered after it has timestamps from all groups in m 's destinations dst . When q reads a timestamp x for m from some group in dst , q updates its clock such that it contains the maximum between its current value and x . Since the next event that q handles for a message m'' will increment its clock, it follows that m'' will have a timestamp bigger than x .

We now show that every message that contains a timestamp smaller than m 's final timestamp ts is eventually delivered and its state set to DONE. To see why, let m' be the message with the smallest timestamp in the buffer. Thus, such a message is eventually delivered and its state set to ORDERED. Eventually, m will be the message in the buffer with smallest timestamp and therefore delivered, a contradiction. We conclude then that q eventually delivers m . \square

PROPOSITION 2. (*Validity*) *If a correct client c multicasts a message m , then eventually every correct process p in m 's destination dst delivers m .*

PROOF: Upon multicasting m , c relays m to groups in dst (see Task 1). The Relay procedure then copies m to the M buffer of every correct process p in groups in dst and sets its state to MCAST. From Lemma 1, it follows that every correct process p set m 's state to ORDERED. From Lemma 2, p eventually delivers m . \square

PROPOSITION 3. (*Uniform agreement*) *If a process p delivers a message m , then eventually all correct processes q in m 's destination dst deliver m .*

PROOF: For process p to deliver m , from Task 6, p has a timestamp for every group h in dst in the M buffer such that ts is the largest among these timestamps. Moreover, there is no message m' in the buffer such that $m \neq m'$, $ts < y$, where y is a timestamp assigned to m' , and m' is not ordered.

We first show by contradiction that q eventually has m in its M buffer. Let c be the client that multicasts m . If c is correct then, c writes m in q 's buffer, so consider that c fails before it can write m in q 's buffer. Since p delivers m , it has a quorum of acknowledgements from each group in dst . Any quorum includes at least one correct process, which from

Task 10, eventually suspects c and relays m to all processes in dst , including q , a contradiction.

It follows from Lemma 1 that q eventually sets the state of m to ORDERED in its buffer, and from Lemma 2 that q eventually delivers m . \square

PROPOSITION 4. (Uniform prefix order) *For any two messages m and m' and any two processes p and q such that $\{p, q\} \subseteq dst \cap dst'$, where dst and dst' are the groups addressed by m and m' , respectively, if p delivers m and q delivers m' , then either p delivers m' before m or q delivers m before m' .*

PROOF: The proposition trivially holds if p and q are in the same group, so assume p is in group g and q is in group h and suppose, by way of contradiction, that p does not deliver m' before m nor does q deliver m before m' . Without loss of generality, suppose that m 's timestamp ts is smaller than m' 's timestamp ts' .

We claim that q inserts m into the M buffer before delivering m' . In order for m (respectively, m') to be delivered by p (resp., q), p 's (resp., q 's) M buffer must contain a timestamp ts_g from group g and ts_h from group h (resp., ts'_g from group g and ts'_h from group h).

From Task 2 (or Task 9 if some process has suspected the leader), the leader l in group g must have included the timestamp ts_g for message m and ts'_g for message m' in p 's

M buffer and both timestamps have been acknowledged by a quorum of processes in group g . Assume that the leader l has written ts_g before ts'_g to the M buffer of every follower in group g and the leader l_h in group h . From Task 2, we have $ts_g < ts'_g$. Therefore, from Task 4, l_h will write to the M buffer of every follower in group h , including q , both ts_g for message m and ts'_g for message m' .

Consequently, from the claim, q delivers m before m' since $m.ts < m'.ts$, a contradiction that concludes the proof. \square

PROPOSITION 5. (Uniform acyclic order) *Let relation $<$ be defined such that $m < m'$ iff there exists a process that delivers m before m' . The relation $<$ is acyclic.*

PROOF: Suppose, by way of contradiction, that there exist messages m_1, \dots, m_k such that $m_1 < m_2 < \dots < m_k < m_1$. From Task 6, processes deliver messages following the order of their final timestamps. Thus, there must be processes p and q such that the final timestamps they assign to m_1 , ts_p and ts_q , satisfy $ts_p < ts_q$, a contradiction since both p and q have the same timestamps for each group in dst in Task 6. \square

THEOREM 1. *RamCast implements atomic multicast.*

PROOF: This follows directly from Propositions 1 through 5. \square